

Practical privacy metrics for synthetic data

Gillian M Raab, Beata Nowok & Chris Dibben

May 14, 2024

Abstract

This paper explains how `synthpop` package¹² has been extended to include functions to calculate identity and attribute disclosure risk measures for the original records, using information obtained from the corresponding synthetic data. The basic function `disclosure` calculates identity disclosure for a set of quasi-identifiers (keys) and attribute disclosure for one variable specified as a target from the same set of keys. The second function `disclosure.summary` is a wrapper for the first and it presents summary results for a set of targets in the data set. This short paper explains the measures of disclosure risk and documents how they are calculated. We recommend two measures: *RepU* (replicated uniques) for identity disclosure and *DiSCO* (disclosive in synthetic and correct in the original) for attribute disclosure. Both are expressed as a % of the original records and both can be compared to similar measures calculated from the original data. Experience with using the functions on real data found that some apparent disclosures could be identified as coming from relationships in the data that would be expected to be known to anyone familiar with its features. We flag cases when this seems to have occurred and provide means of excluding them.

1 Introduction

In his recent review of thirty years of synthetic data Reiter [16] comments:

"While there is need to examine disclosure risks in synthetic data, there is no standard for doing so, especially in fully synthetic data. Instead, disclosure checks tend to be ad hoc"

¹see www.synthpop.org.uk

²The new version 1.8-1 is currently available on Github at <https://github.com/gillian-raab/synthpop>

This is in contrast to the variety of measures of utility available for synthetic data; see [15] for a list of these. While utility measures must be chosen that are relevant to the intended use of the synthetic data, disclosure measures must focus on the possible harm to the privacy of an individual or other unit whose data contributed to the creation of the synthetic data. The evaluation of the disclosure risk from synthetic data must relate to the context of its release (see [4] for a discussion of this). We cannot expect a fixed rule, for example that a criterion for release requires a value of some disclosure measure beyond a threshold. Instead, we expect those releasing data to use the disclosure measures to evaluate potential harm to the data subjects or to the data custodians from information in the released data. We hope that these new functions will allow the disclosure risk of synthetic data sets to be explored and, where necessary, reduced.

We expect that disclosure risks from synthetic data to be lower than those from the original data. But in some cases, e.g. a large data set with only a few categorical variables, the disclosure risk from the original may already be low. To evaluate the disclosure risk of a synthesis method, the risks for synthetic data must be compared with equivalent risks for the original. We consider two types of disclosure risk:

- *identity disclosure*: This refers to the ability to identify individuals in the data from a set of known characteristics that we will refer to as keys. Identity disclosure may be less relevant for completely synthesized³ data because there is no one-to-one correspondence between records in the original and synthetic data sets. But it may still be of interest since it is an important factor for attribute disclosure.
- *attribute disclosure*: This refers to the ability to find out from the keys something, not previously known, for an attribute associated with a record in the original data.

The disclosure risk posed by synthetic data can be reduced by using techniques from statistical disclosure control (**sd**c), such as aggregation of categories, smoothing of numeric values or removal of replicated uniques. These methods can be used to reduce disclosure risk by modifying the original data before synthesis, or the synthetic data before its is released. Some such methods are already available in the **synthpop** package. These include categorising, top/bottom coding and smoothing for continuous variables, and the merging of small categories for factors. The removal of replicated uniques is another option available.

³Complete or full synthesis is when all values of all variables are replaced by synthetic values. This is in contrast to incomplete or partial synthesis where only some variables are replaced.

There has been many recent proposals for making synthetic data sets Differential Private (DP) (selection of refs). DP is a very strong privacy guarantee that protects against an intruder with arbitrary external information about the subjects in the data, except for the one whose privacy is being protected. This is an unrealistic assumption and DP synthetic data has been shown to have unacceptably low utility in many cases [1, 18]. We will not discuss these methods here, but note that we could use the metrics proposed here to evaluate disclosure risks for DP synthetic data.

2 A simple example

Here we illustrate the basic use of `disclosure.summary`. If the parameter `targets` is not specified, all the variables in the synthetic data that are not part of keys are used as targets. The identity disclosure measures are `Ui0` for original and `repU` for synthetic, and for attribute disclosure `Di0` for original and `DiSC0` for synthetic. These and other measures will be explained in Sections 3.2 and 3.3.

First, a subset of 9 variables are selected from the `SD2011` data (a survey on quality of life in Poland) that is available as part of the `synthpop` package. A single synthetic data set is created by the default method in `synthpop`: `cart` for each conditional distribution. The synthetic data object `s1`⁴ has a component `syn` that is a single synthetic data set. The disclosure functions can also be used with synthetic data created by other methods either as single synthetic data sets or lists of repeated syntheses from the same original. Here we select 4 keys that represent items that might be known for members of this sample, or of the Polish population in 2011. By default, the 5 other variables become the targets: `income` (rounded to give 406 distinct values) `marital` (marital status), `workab` (intention to work abroad), `ls` and `depress` (scores for life satisfaction and depression with 7 and 21 distinct values). We first use `disclosure.summary` to create a `disclosure.summary` object for these keys, and we print out its identity component.

```
R> library(synthpop)
R> ods <- SD2011[, c("sex", "age", "region", "placesize", "depress",
+   "income", "ls", "marital", "workab")]
R> s1 <- syn(ods, seed = 8564, print.flag = FALSE)
R> t1 <- disclosure.summary(s1, ods, print.flag = FALSE, plot = FALSE,
+   keys = c("sex", "age", "region", "placesize"))
R> print(t1, to.print = "ident")
```

⁴an object of class `synds`

Disclosure risk for 5000 records in the original data

```
Identity disclosure measures
from keys: sex age region placesize
For original ( UiO ) 48.38 %
For synthetic ( repU ) 14.86 %.
```

The measure *UiO* (Unique in Original) shows that 48% of the original records would have unique combinations of these 4 keys. The term, "singling out" is used in data protection regulation for this type of attribute disclosure⁵. For the synthetic data *RepU* tells us that almost 15% of the original records would be unique in the original and also unique in the synthetic data.

For attribute disclosure we can examine the results as either a table or a plot. Here we see in Figure 1 the plot that would have been generated if `plot` had been set to `TRUE` in the code above, the attribute disclosure results would have been plotted as shown in Figure 1. Details of how identity and disclosure measures are calculated can be found in the Section 3.




Figure 1: Plot from `disclosure.summary` with `plot = TRUE`, the default value.

The results in the plot are in descending order by the disclosure in the synthetic data.

The measure *DiO* tells us that these 4 keys would identify a unique value of each of these targets for over 50% of the original records. *DiSCO*, the proportion of the original records that are disclosive in the original and also in the synthetic data with a correct attribution to the target give lower values. The first two variables shown in Figure 1 have additional labels that flag possible

⁵See for example its use in the UK Information Commissioner's guidance on anonymisation here <https://ico.org.uk/media/about-the-ico/documents/4018606/chapter-2-anonymisation-draft.pdf>

contributions to disclosure from knowledge of 1-way or 2-way relationships in the original data. This is detailed in Section 4.

Previous work on attribute disclosure [3, 20] has used the Correct Attribution Probability (*DCAP*) as a disclosure measure for synthetic data. This is calculated as the percentage of records correctly predicted from the synthetic data by a set of keys. However it has been suggested [2, 6], that it may be as much a measure of utility as disclosure risk and modifications have been suggested [9, 19] one of which (*TCAP*) is close to our proposed measure, *DiSCO*. For completeness the measures *DCAP* and *TCAP* are calculated by `disclosure`; see Appendix 2.

3 Scenario and definitions

3.1 Setting the scene

These disclosure measures are intended to assess what a person who only has access to the synthetic data can infer about known individuals who are present in the original data. We will use the term "intruder" for such a person, though no malicious intent is implied. The intruder is assumed to have information for one or more individuals about the value of certain key variables that are present in the same format in the original and the synthetic data. They attempt first to see if the individual is present, and then to determine the value of other items in the data file that we refer to as targets. We are assuming a worst-case scenario where the intruder believes they are querying the original data.⁶ Disclosure measures from the synthetic data are each compared to similar measures for someone with access to the original data. Here we will introduce the measures by an example. Formal definitions with notation and formulae are in Appendix 1. The first step in evaluating disclosure risk, as described here, is to identify a set of keys that might be expected to be known to an intruder. These keys are then combined to form a quasi-identifier that we designate as q . For example, if we have hospital records we might define age, sex date and hospital as keys and this would give a q with levels such as "78 | M | 1/1/2024 | WG" for a 78 year old man admitted to hospital WG on 1/1/2024.

⁶This may not be too unrealistic if the data are made available inadvertently, or if the intruder thinks that efforts to label the synthetic data as e.g. "Fake Data" are thought just to be a cover up.

3.2 Identity disclosure measures

The concept of k -anonymity is central to identity disclosure for microdata. First proposed in 1998 [17] it is discussed fully in [4]. A table is k -anonymous if a set of keys identifies at most $k - 1$ individuals. Thus 2-anonymous data will never identify just one individual. Based on this idea, the percentage of records for which the keys identify just one individual give identity disclosure measures. A table of q values is produced from the synthetic and the original data. UiO and UiS are the percentages of records with keys where the table value is 1. An intruder checking out a record for their known set of keys will look for it in the synthetic data. Some records will not be in the synthetic data and $UiOiS$ (Unique in Original in Synthetic) gives the % that would be found. These records are then checked for uniqueness in the synthetic data giving, $repU$ is the percentage of unique original records that are also unique in the synthetic data.

The percentage $repU$ has been used as a disclosure measure to evaluate synthetic data by [8] and by [14]⁷. Replicated uniques are used in `synthpop` as part of the statistical disclosure control function, `sdc`, that includes the option of reducing disclosure risk by removing them from the synthetic data. Nowok et al. [12] have evaluated this and give an example where this process has very little effect on utility. The function `replicated.uniques`⁸ also calculates $repU$ using a different method from the one described here.

One of the outputs of the functions `disclosure` and `disclosure.summary` is `ident`, a table of identity disclosure measures as illustrated by this example, using the keys from the example in Section 2 but now calculated for a synthetic object with 5 data sets and using the one target `depress`. We first create `t5` an object of class `disclosure` and then print out the identity and attribute disclosure measures for each synthetic data set.

The table for identity disclosure has DiO and $repU$ as its first and last column. The 2nd and 3rd columns are UiS calculated from the synthetic data in the same manner as UiO from the original and $UiOiS$ the % of UiO with q that are in the synthetic data but not necessarily unique. Each are steps towards calculating $repU$.

```
R> s5 <- syn(ods, seed = 8564, m = 5, print.flag = FALSE)
R> t5 <- disclosure( s5, ods, keys = c("sex", "age", "region",
+   "placesize"), target = "depress", print.flag = FALSE)
R> print(t5, to.print = c("ident"))
```

⁷Jackson et al. in [8] argue that the denominator for $repU$ should be N_s rather than N_d . This is inappropriate because our scenario is to consider the risk to the original data.

⁸For example by `replicated.uniques (s2, ods, keys = c("sex","region","age","placesize")`

Disclosure measures from synthesis for 5000 records in original data.

Identity disclosure measures for 5 synthetic data set(s) from keys:
sex age region placesize

	UiO	UiS	UiOiS	repU
1	48.38	37.34	22.68	14.86
2	48.38	35.44	22.24	13.96
3	48.38	35.18	21.98	13.62
4	48.38	34.90	22.08	13.78
5	48.38	36.14	22.00	14.62

3.3 Attribute disclosure measures

We approach disclosure risk from the point of view of an intruder with access to the synthetic data and to certain attributes (quasi-identifiers) known for one or more individuals in the original data. The quasi-identifiers for the known record(s) can be combined to create a composite variable q that can be created for all records in the original and synthetic data. The intruder will proceed as follows:

- Look up q for each original record in the synthetic if no records are found score 0.
- For those found (in synthetic original, iSO as a proportion of original records), check if all such records have the same level of the target t in the synthetic data as in the original. Score 1 for each such original record where this apparent disclosure corresponds to the value of t in the original data; Otherwise score 0. The mean of this score is $DiSCO$, disclosive in synthetic correct in original.

Note that records contributing to $DiSCO$ may not be disclosive in the original data, as this information would not be available to the intruder. A further measure $DiSDiO$ (Disclosive in Synthetic Disclosive in Original) restricts the score to those disclosive in the original. These measures are defined formally in Appendix 1.

```
R> print(t5, to.print = c("attrib"))
```

Disclosure measures from synthesis for 5000 records in original data.

	DiO	DiS	iSO	DiSCO	DiSDiO	max_denom	mean_denom
1	53.3	46.26	64.90	9.54	6.14	3	1.16

2	53.3	44.80	64.00	10.26	6.78	4	1.19
3	53.3	44.60	64.02	9.10	5.92	4	1.19
4	53.3	45.80	63.88	9.20	5.52	4	1.21
5	53.3	44.52	63.44	9.34	5.80	4	1.23

Here we print the table of attribute disclosure measures for the example in the previous section. As we move from *DOS* to *DiSCO* and then *DisDiO* we can see how the %disclosive is affected by different conditions. Here lack of q levels in the synthetic data reduces disclosure to 64%; then the requirement for correct attribution reduces it to around 9%; finally restricting to records disclosive in the original reduces it to 6%,

The *DiO* and *DiSCO* measures are not restricted to disclosures that are identified from unique records for q in either the original or the synthetic data. The number of records contributing to each disclosive qt cell in the synthetic table is the denominator that applies to that record. The columns `max denom` and `mean denom` refer to the denominators for the records disclosive in the synthetic data that contribute to the *DiSCO* or *DISCDO* measures. We can see from the mean that here the majority of disclosive records had unique key combinations in the synthetic data, and the maxima was 4, 5 or 6 for different syntheses. Large denominators can be an indication that some of the disclosure may be coming from strong relationships between variables in the data that might even be expected a-priori. This aspect is discussed further in Section 4. The disclosure measures can be restricted to those with small denominators by using the parameter `exclude_over_denom_lim` to `TRUE`. The example in Section 2 is here run restricted to denominators of 1.

```
R> disclosure.summary(s1, ods, print.flag = FALSE, plot = FALSE,
+   keys = c("sex", "age", "region", "placesize"),
+   denom_lim = 1, exclude_ov_denom_lim = TRUE)
```

Disclosure risk for 5000 records in the original data

```
Identity disclosure measures
from keys: sex age region placesize
For original ( Ui0 ) 50.26 %
For synthetic ( repU ) 16.36 %.
```

```
Table of attribute disclosure measures for sex age region placesize
Original measure is DiO and synthetic measure is DiSCO
Variables Ordered by synthetic disclosure measure
```

```
attrib.orig attrib.syn check1 Npairs check2
```


1	income	48.38	4.02	0
2	depress	48.38	7.02	0
3	ls	48.38	9.74	0
4	marital	48.38	15.72	0
5	workab	48.38	20.62	0

The *DiSCO* values have decreased, as expected, but *UiO* and *repU* have increased and also that *DiO* is now the same for all targets and equal to the *UiO* value before the denominator exclusion. This makes sense because removing large denominators from *UiO* and *repU* increases the number of uniques. Also, with denominators of 1, all unique values of q are disclosive in the original data. for replicated uniques. Note that decreases in *DiSCO* are more pronounced for *workab* and *marital*, the two targets that were flagged to check 1-way or 2-way relationships. Their original *DiSCO* values were 37% for *marital* and 53% for *workab*. This and other methods of excluding certain apparent disclosures are discussed in the next section.

4 Identifying disclosure from 1-way and 2-way relationships

As mentioned in the Introduction, what we can learn about disclosiveness of attributes can depend on our prior knowledge of the data set or the population from which it is drawn. It would not be practical to specify our prior probability for every possible combination of keys. However, checking two aspects of disclosure results can help us to check when we might have predicted the correct attribution with high probability without knowledge of all the keys. The first is a check for a target where a high proportion of records have one level of the target. The second is when there is a strong relationship between the target and one of the keys, so that one tq pair accounts for many of the disclosive records. These two aspects are flagged by the values of `check_1way` and `check_2way` that are returned as part of objects of `Returning` to the example in 2 we now use the function `disclosure` to get details for the two targets that were flagged as requiring checking; see Figure 1.

```
R> d1_workab <- disclosure(s1, ods, print.flag = FALSE, target = "workab",
+   keys = c("sex", "age", "region", "placesize"), plot = FALSE)
R> print(d1_workab, to.print = c("check_1way", "check_2way"))
```

Disclosure measures from synthesis for 5000 records in original data.

```
Details of target level contributing disproportionately to disclosure
check Level totalDisclosive PctDisAll totalLevel pctdisLevel
```

1	1	NO	4430	88.6	4116	92.91
---	---	----	------	------	------	-------

Details of target-key combinations contributing disproportionately to disclosure
This is a list with one for each of syntheses

```
[[1]]
  total pctall max_denom   key   target_key pct_2way denom_2way
1     6    0.2         6   age          NO | 55    97.1      105
2     6    0.2         6 region NO | Podkarpackie  95.2      313
```

We can see that it was the category "NO" of `workab` that contributed most to the disclosure risk; most survey respondents (93%) had never worked abroad. This level of the target accounted for 89% of apparent disclosures for `workab`. To predict this level for all of a group with the same q could hardly be considered disclosive. This would be true if the intruder had access to the marginal distribution of `workab`, and even if they did not some knowledge of the respondents' background might suffice. The tq pairs identified for `workab` both included the "NO" level of the target.

```
R> d1_marital <- disclosure(s1, ods, print.flag = FALSE, target = "marital",
+   keys = c("sex", "age", "region", "placesize"), plot = FALSE)
R> print(d1_marital, to.print = c("check_1way", "check_2way"))
```

Disclosure measures from synthesis for 5000 records in original data.
No 1 way checks for your target are flagged at your settings for `thresh_1way`

Details of target-key combinations contributing disproportionately to disclosure
This is a list with one for each of syntheses

```
[[1]]
  total pctall max_denom key   target_key pct_2way denom_2way
1     7    0.4         7 age MARRIED | 40    82.3      79
2     6    0.3         6 age  SINGLE | 20   100.0      79
```

For the target `marital` two tq pairs are identified as contributing large denominators. Those aged 20 are all SINGLE, and those aged 40 are mainly MARRIED. Again these would not be considered disclosive for an intruder with some knowledge of the respondents.

The thresholds for identifying one-way and two-way relationships can be modified. In each case there are two criteria that need to be satisfied, one number and one %. For one-way disclosure the parameter `thresh_1way` has the default value of $c(50, 90)$, meaning that there must be at least 50 disclosive records for one level of a target and that q values including this target must account for over 90% of all disclosive records. For two way relationships the `thresh_2way` has default value $c(5, 80)$ and the algorithm first identifies all disclosive tq combinations with denominators over 5.

It then identifies the level of the key in q with that best predicts this level of t and checks if over 80% of the disclosive records with this key would have the correct prediction.

5 Excluding records

Having identified records where the disclosure is not really telling us anything new, one option is to exclude some of these original records explicitly from the measures. The following parameters for `disclosure.summary` and `disclosure` can be used for this.

- **not.target**: All records with levels of the target given by the parameter **not.target** are excluded from all disclosure measures.
- **usekeysNA**: This is set to TRUE by default so that missing values are included in all tables. It can be set to FALSE for some or all keys to exclude NAs.
- **usetargetsNA(disclosure.summary)** or **usetargetNA(disclosure)**: Similar to the above for **target(s)** - note can be a vector for `disclosure.summary`.
- **exclude.keys**, **exclude.keylevs** and **exclude.targetlev**: For disclosure only: Three vectors of length the number of key-target pairs to be excluded from all tables.

To illustrate exclusions we will use the Adult data set from the UCI machine learning repository [11] with almost 50 thousand records from the US Census income study, available as part of the `arules` package for **R** [7]. This is one of the data sets used in [5] to evaluate privacy risks for synthetic data. Figure 2 is the output of `disclosure.summary` from the keys `age`, `occupation`, `race` and `sex` for the other 10 variables in the data. The code used to calculate the results in this section is in Appendix 3. The disclosiveness of the original data is relatively low, compared to the previous example, because of the large sample size and the absence of any geographic identifiers. Three of the 10 variables `capital.gain`, `capital.loss` and `native.country` are flagged to check one-way relationships. The `check1` column of the attribute table tells us that the levels contributing to disclosure are zeros for the first two and `United-States` for `native.country`. These levels make up 95%, 92% and 89% of all records.

Three other variables are flagged as having disclosive two_way relationships `workclass`, `marital`, `relationship` with totals of 7, 7 and 2 tq combinations respectively. The summary




Figure 2: Plot from `disclosure.summary` for variables flagged by checks.

function `disclosure.summary` does not allow target and key specific pairs to be excluded. To investigate this it is necessary to examine the output of `disclosure`⁹. This showed that the largest contribution to `check_2way` for `workclass` was due to this being missing when occupation was missing, although other relationships between these two variables also contributed.

Table 1 gives the results of excluding different entries in the tables of q and t from the attribute disclosure measures. Excluding the levels of the target flagged by `textttcheck_1way` reduces the *DiSCO* to almost zero for these 3 variables. Adding exclusion of missing values reduced the disclosure for `workclass` and for some other variables a little. Adding the restriction to denominators of 1, reduced the disclosure for variables identified by `check_2way` to low levels. In the final columns we can see that restricting to denominators of 1, by itself, gives low levels of disclosure for all variables, with the exception of those flagged by `check_1way`. This approach was one that was trialed in an earlier version of these functions but discarded for giving misleading results for some data. It does not count some attribute disclosures that might well be found by an intruder. For example, if a small number of records with the same keys all had the same level

⁹see code in Appendix 3 for how to do this.

target	No excl		not.tlev		NAs out not.tlev		denom_lim 1 NAs out not.tlev		denom_lim 1	
	orig	syn	orig	syn	orig	syn	orig	syn	orig	syn
capital.gain	22.55	19.87	0.21	0.00	0.21	0.00	0.19	0.00	2.68	1.17
capital.loss	30.61	27.71	0.08	0.00	0.08	0.00	0.08	0.00	2.68	1.26
education	3.71	0.93	3.71	0.93	3.71	0.93	2.45	0.46	2.68	0.50
fnlwgt	2.70	0.00	2.70	0.00	2.70	0.00	2.45	0.00	2.68	0.00
hours.per.week	4.36	0.19	4.36	0.19	4.36	0.19	2.45	0.12	2.68	0.13
income	4.97	2.10	4.97	2.10	4.97	2.10	1.58	0.50	2.68	0.82
marital.status	8.23	5.27	8.23	5.27	8.23	5.27	2.45	0.68	2.68	0.74
native.country	17.09	13.55	0.94	0.08	0.94	0.08	0.67	0.04	2.68	0.87
relationship	5.17	2.64	5.17	2.64	5.17	2.64	2.45	0.69	2.68	0.73
workclass	14.27	11.64	14.27	11.64	14.27	11.64	2.45	0.85	2.68	0.96

Table 1: Disclosure results from Adult data with different exclusions.

of the target that corresponded to the level in the original data. We feel that a better approach is to exclude specific target/key combinations.

6 Conclusions

The privacy metrics we propose here are in some sense the opposite of differential privacy (DP). DP claims to protect data from an intruder with arbitrary knowledge of the data, except of the one record that has the greatest influence on the likelihood of the results. In contrast our metrics require the user to specify the details of what they would expect an intruder to know about the data. Most importantly, the user must specify keys that identify variables in the data that would expect to be known about individuals. Also, the routines flag cases where part of the disclosure measures come from one- or two-way relationships so that disclosure would be expected even in the absence of data. We have evaluated the routines on a few data sets and set levels of the thresholds for this at what seem to be reasonable levels, but more experience with other data sets would be valuable.

We hope that these tools will be helpful to data holders who need to make decisions about the risks of releasing synthetic data to the public, or to a restricted audience. They should also enable the disclosure risks of different synthesis methods to be evaluated.

7 Acknowledgement

Research Data Scotland (<https://www.researchdata.scot/>) funded Gillian Raab’s time to carry out the research reported here and to expand the capabilities of the synthpop package¹⁰ to include measures of disclosure control. We also thank the Scottish Centre for Administrative Data Research for continue to support the development of the synthpop package since its creation was supported by the ESRC funded SYLLS project in 2012-14.

References

- [1] BOWEN, C. M., AND LIU, F. Comparative study of differentially private data synthesis methods. *Statistical Science* 35 (2020), 280—307.
- [2] CHEN, TAUB, J., AND ELLIOT, M. Trade-off between information utility and disclosure risk in ga synthetic data generator. UNECE Work Session on Statistical Data Confidentiality, Skopje, North Macedonia. Available from https://unece.org/fileadmin/DAM/stats/documents/ece/ces/ge.46/2019/mtg1/SDC2019_S3_UK_Chen_Taub_Elliot_AD.pdf, 2019. Accessed: 2024-01-06.
- [3] ELLIOT, M. Final report on the disclosure risk associated with the synthetic data, produced by the sylls team. Available from <https://tinyurl.com/syllsDR>, 2014. Accessed: 2022-02-23.
- [4] ELLIOT, M., MACKEY, E., AND O’HARA, K. The anonymisation decision-making framework: European practitioners. Available from <https://ukanon.net/framework/>, 2020. Accessed: 2022-02-23.
- [5] GIOMI, M., BOENISCH, F., WEHMEYER, C., AND TASNÁDI, B. A unified framework for quantifying privacy risk in synthetic data, 2022.
- [6] GM, R. Comments on disclosure measure used in data challenge report. Letter send to contributors to synthetic data Challenge at Newton Institute from, 2018.

¹⁰from version 1.8-0 available on CRAN at <https://CRAN.R-project.org/package=synthpop>

- [7] HAHLER, M., GRUEN, B., AND HORNIK, K. arules – A computational environment for mining association rules and frequent item sets. *Journal of Statistical Software* 14, 15 (October 2005), 1–25.
- [8] JACKSON, J., MITRA, R., FRANCIS, B., AND DOVE, I. Using saturated count models for user-friendly synthesis of large confidential administrative databases. *Journal of the Royal Statistical Society: Series A (Statistics in Society)* 185 (2022), 1613—1643.
- [9] LITTLE, C., ELLIOT, M., AND ALLMENDINGER, R. Comparing the utility and disclosure risk of synthetic data with samples of microdata. In *Privacy in Statistical Databases* (Cham, 2022), J. Domingo-Ferrer and M. Laurent, Eds., Springer International Publishing, pp. 234–249.
- [10] MACHANAVAJJHALA, A., GEHRKE, J., KIFER, K., AND VENKITASUBRAMANIAM, M. l-diversity: Privacy beyond k-anonymity. In *22nd International Conference on Data Engineering (ICDE’06)* (2006), IEEE.
- [11] NEWMAN, C. B. D., AND MERZ, C. UCI repository of machine learning databases, 1998.
- [12] NOWOK, B., RAAB, G., AND DIBBEN, C. Recognising real people in synthetic microdata: risk mitigation and impact on utility. UNECE Work Session on Statistical Data Confidentiality, Skopje, North Macedonia. Available from https://unece.org/fileadmin/DAM/stats/documents/ece/ces/ge.46/2017/3_risk_mitigation.pdf, 2017. Accessed: 2022-02-23.
- [13] PATER, L., AND SMID, S. Making attribute information of synthetic data interpretable with the aggregation equivalence level. UNITED NATIONS ECONOMIC COMMISSION FOR EUROPE CONFERENCE OF EUROPEAN STATISTICIANS Expert Meeting on Statistical Data Confidentiality 26-28 September 2023, Wiesbaden. Available from https://unece.org/sites/default/files/2023-08/SDC2023_S4_3_Netherlands_Pater_D.pdf, 2024. Accessed: 2022-02-23.
- [14] RAAB, G. Utility and disclosure risk for differentially private synthetic categorical data. In *Privacy in Statistical Databases, 2022* (2022), K. Muralidhar and J. Domingo-Ferrer, Eds., Springer, Berlin.

- [15] RAAB, G. M., NOWOK, B., AND DIBBEN, C. Assessing, visualizing and improving the utility of synthetic data, 2021.
- [16] REITER, J. Synthetic data: A look back and a look forward. *Transactions in Data Privacy* 16 (2023), 15—24.
- [17] SAMARATI, P., AND SWEENEY., L. Generalizing data to provide anonymity when disclosing information. In *Proceedings of the Seventeenth ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems* (1998), ACM Press.
- [18] STADLER, T., OPRISANU, B., AND TRONCOSO, C. Synthetic data – anonymisation groundhog day. In *31st USENIX Security Symposium (USENIX Security 22)* (Boston, MA, 2022), pp. 1451–1468.
- [19] TAUB, J., AND ELLIOT, M. The synthetic data challenge. UNECE Work Session on Statistical Data Confidentiality, Skopje, North Macedonia. Available from <https://unece.org/fileadmin/DAM/stats/documents/ece/ces/ge.46/2019/mtg1/SDC2019>, 2019. Accessed: 2024-01-06.
- [20] TAUB, J., ELLIOT, M., TAUB, M., AND SMITH, D. Differential correct attribution probability for synthetic data: An exploration. In *Privacy in Statistical Databases - UNESCO Chair in Data Privacy, International Conference, PSD 2018, Proceedings* (2018), F. Montes and J. Domingo-Ferrer, Eds., Springer-Verlag Italia, pp. 122–137.

Appendix 1: Notation and formal definitions

Before defining the measures of identity and disclosure risk we need to introduce the notation that will be used to calculate them. The first step is to create the quasi-identifiers from the keys for the original and synthetic data. For the keys used in the example given in Section 2 the quasi-identifier that we will designate as q for the first record in the original data is:

"FEMALE | 57 | Lubuskie | URBAN 100,000-200,000"

and that for the first record in the synthetic data:

"FEMALE | 39 | Zachodnio-pomorskie | URBAN 100,000-200,000".

In order to calculate identity disclosure measures, we need to compare the tables of q from the original and synthetic data. For attribute disclosure measures we need to cross-tabulate q with each target variable t and compare findings from the synthetic data with what would have been found from the original data. In general, the levels of q and sometimes t in the original and synthetic data will not be the same. Before creating any tables, we need to define sets of q and t values that give the union of both sets of levels and align the tables so that their indices correspond.

For the original data $d_{.q}$ is the count of records with the keys corresponding to the levels of q and q_{tq} the count of records with this q and level $t = 1, \dots, T$ of the target. The equivalent counts from the synthesised data are designated by $s_{.q}$ and s_{tq} . When a member of q is in the original data but not in the synthetic, $s_{.q}$ and s_{tq} are all zero. Similarly when a member of q is in the synthetic data but not in the original, $d_{.q}$ and d_{tq} are all zero. The two tables can be written as shown in Table 1, where the total records in the original data is N_d , made up of $N_{d \text{ only}}$ and $N_{d \text{ both}}$. The equivalent totals for the synthetic data are N_s , $N_{s \text{ only}}$ and $N_{s \text{ both}}$.

	only in original			in both			only in synthetic			Total
1	...	d_{1q}	d_{1q}	0	...	$d_{1.}$
...
t	...	d_{tq}	d_{tq}	0	...	$d_{t.}$
...
T	...	d_{Tq}	d_{Tq}	0	...	$d_{T.}$
Column sums	$d_{.q}$...	$d_{.q}$	0	...	N_d
Totals	$N_{d \text{ only}}$			$N_{d \text{ both}}$			0			N_d

	only in original			in both			only in synthetic			Total
1	...	0	s_{1q}	s_{1q}	...	$s_{1.}$
...
t	...	0	s_{tq}	s_{tq}	...	$s_{t.}$
...
T	...	0	s_{Tq}	s_{Tq}	...	$s_{T.}$
Column sums	...	0	$s_{.q}$	$s_{.q}$...	N_s
Totals	0			$N_{s \text{ both}}$			$N_{s \text{ only}}$			N_s

Table 2: Notation for tables from quasi-identifier (q) and target (t) from original (upper table) and synthetic data (lower table).

To calculate the % of records in the original and synthetic data we need:

$$\% \text{ Unique in Original} = UiO = 100 \sum (d_{.q} | d_{.q} = 1) / N_d. \quad (1)$$

$$\% \text{ Unique in Synthetic} = UiS = 100 \sum (s_{.q} | d_{.q} = 1) / N_d. \quad (2)$$

The intruder has information about the keys for an individual in the real data that they attempt to identify in the synthetic data. They first attempt to find them in the synthetic data, and the % found is:

$$\% \text{ Unique in Original in Synthetic} = UiOiS = 100 \sum (d_{.q} = 1 | s_{.q} = 1 \wedge d_{.q} > 0) / N_d. \quad (3)$$

Some of these records would not be unique in the synthetic data, restricting to such records gives:

$$\% \text{ replicated Uniques} = repU = 100 \sum (s_{.q} | d_{.q} = 1 \wedge s_{.q} = 1) / N_d. \quad (4)$$

To find an attribute from a set of keys, it is necessary to examine the distribution of s_{tq} for groups defined by q . We define column proportions for the original and synthetic data as $pd_{tq} = d_{tq} / d_{.q}$ and for the synthetic as $ps_{tq} = s_{tq} / s_{.q}$.

Returning to the scenario described in Section 3.1, we must first define a measure of attribute disclosure for the original data. This is based on the concept of *l-diversity* [10] that requires that each set of records defined by q has at least $l (\geq 2)$ distinct values of the target. A data set is *l2-diverse* for q and t if all records for every q have the same level of t ¹¹. An attribute disclosure measure for the original data can be defined as *% Disclosive in Original* :

$$DiO = 100 \sum_q \sum_t (d_{tq} | pd_{tq} = 1) / N_d. \quad (5)$$

An intruder with access only to the synthetic data, but with knowledge of q from one or more individuals in the original, would look them up in the synthetic data. Some of their q levels be key combinations that do not appear in the synthetic data leaving the proportion that do appear

¹¹This could be generalised to $l > 2$, but in practice the levels of targets are not generally exchangeable and a more practical approach would be to aggregate levels for certain targets.

as *iSO* (in in Synthetic Original)

$$iSO = 100 \sum_{q=1}^q \sum_{t=1}^t (d_{tq} | s_{tq} > 0) / N_d. \quad (6)$$

A level of q from an original record may identify more than one target in the synthetic data, or identify the wrong target. To exclude these we require the records to be Disclosive in the Synthetic data and Correct when checked with the original giving:

$$DiSCO = 100 \sum_{q=1}^q \sum_{t=1}^t (d_{tq} | p_{s_{tq}} = 1) / N_d. \quad (7)$$

Note that *DiSCO* can include records that are not disclosive in the original data giving a further measure Disclosive in Synthetic and Disclosive in the Original :

$$DiSDiO = 100 \sum_{q=1}^q \sum_{t=1}^t (d_{tq} | p_{s_{tq}} = 1 \wedge p_{d_{tq}} = 1) / N_d. \quad (8)$$

As we comment above the intruder would not be able to tell if records were identified as *DiSDiO* rather than *DiSCO*, so we prefer the latter measure. However, the intruder can identify when the apparently disclosive record is not unique in the synthetic data. This restriction can be imposed by requiring the denominator in the synthetic data to exceed a 1, as described and discussed in Section 5.

Appendix 2: CAP measures

The measures *baseCAPd*, *DCAP*, *TCAP* are calculated by the function `disclosure` and are stored in the component `allCAPs` of the output object of class `disclosure`. This is printed when the parameter `to.print` includes `allCAPs`. The first measure is known as the baseline CAP and refers to an average of the predictions that would be made by someone who only has access to the marginal distribution of the target. The intruder then guesses the CAP for each level of the target according to the relative frequencies pd_t . Averaging this over all observations gives

$$baseCAPd = \sum (pd_t)^2 / N_d.$$

The % of t that would be correctly predicted from q for someone with access to the original data is *CAPd*:

$$CAPd = \sum_{tq} (pd_{tq}d_{tq})/N_d,$$

and the equivalent measure for the synthetic data treated as if it were the original is

$$CAPd = \sum_{tq} (ps_{tq}s_{tq})/N_s.$$

The measure *DCAP* is the percentage of t correctly predicted in the synthetic data q . This gives

$$DCAP = \sum_{tq} (ps_{tq}d_{tq})/N_d$$

The *DCAP* score can be expressed by scaling it from *baseCAPD* to $1[9, 13]$, although this can result in some negative values. The *TCAP* measure as defined by [9] uses the same count of disclosive records as *DISCO*, but uses as denominator the number of records that have keys represented in both the synthetic and the original data. It can also be scaled from *baseCAPD* to 1, and as [9] demonstrate this measure can also take negative values.

```
R> print(t5, to.print = "allCAPs")
```

Disclosure measures from synthesis for 5000 records in original data.

```
CAP measures for the target depress with keys
sex age region placesize
```

	baseCAPd	CAPd	CAPs	DCAP	TCAP
1	9.81	74.15	69.78	16.39	14.70
2	9.81	74.15	69.36	17.45	16.03
3	9.81	74.15	69.24	16.20	14.21
4	9.81	74.15	69.87	15.92	14.40
5	9.81	74.15	69.00	16.17	14.72

As expected the lower denominator for *TCAP* gives values higher than the equivalent *DiSCO* values for the same example given in Section 3.3.

Appendix 3:: Code for analysis of the Adult data set.

```
library(arules)
```

```

data(AdultUCI)
names(AdultUCI) <- gsub("-", ".", names(AdultUCI))
dim(AdultUCI)

##----- drop one of two identical education variables-----
adultsample <- AdultUCI[,-c(5)]
names(adultsample)[-c(1,6,8,9)]
names(synadult$syn)
##----- native.country moved to end as many levels that slow down synthesis-----
synadult <- syn(adultsample, visit.sequence = c(1:12,14,13) ,seed = 7868)
##----- get basic disclosure and make Figure 2-----
dd1 <- disclosure.summary(synadult,adultsample,
  keys = c("age" , "occupation","race", "sex" ))
png("vignettes/fig2dis.png",height=300, width=500)
print(dd1)
dev.off()
##----- exclude some target levels and missing values for targets-----
dd2 <- disclosure.summary(synadult,adultsample,
  print.flag = FALSE, not.targetslev = c(rep("",5),"0","0","", "United-States",""),
  keys = c("age" , "occupation","race", "sex" ))
##----- exclude some target levels and missing values for targets-----
dd3 <- disclosure.summary(synadult,adultsample, print.flag = FALSE,
  not.targetslev = c(rep("",5),"0","0","", "United-States",""),

  keys = c("age" , "occupation","race", "sex" ))
dd1 <- disclosure.summary(synadult,adultsample, keys = c("age" , "occupation","race",
##-----restrict to denominators of 1 -----
dd4 <- disclosure.summary(synadult,adultsample, denom_lim =1,
  not.targetslev = c(rep("",5),"0","0","", "United-States",""),
  usekeysNA = FALSE, usetargetsNA = FALSE,

```

```

    exclude_ov_denom_lim = TRUE,
    keys = c("age" , "occupation","race", "sex" ))
##-----just denominators not others-----
dd5 <- disclosure.summary(synadult,adultsample, denom_lim =1,
    exclude_ov_denom_lim = TRUE,
    keys = c("age" , "occupation","race", "sex" ))
##-----make table 1-----
library(stringr)
cols12 <- dd1$attrib.table[,1:2]
  dimnames(cols12)[[1]] <- word(dimnames(cols12)[[1]],2,2)
cols34 <- dd2$attrib.table[,1:2]
  dimnames(cols34)[[1]] <- word(dimnames(cols34)[[1]],2,2)
cols56 <- dd3$attrib.table[,1:2]
  dimnames(cols56)[[1]] <- word(dimnames(cols56)[[1]],2,2)
cols78 <- dd4$attrib.table[,1:2]
  dimnames(cols78)[[1]] <- word(dimnames(cols78)[[1]],2,2)
cols910 <- dd5$attrib.table[,1:2]
  dimnames(cols910)[[1]] <- word(dimnames(cols910)[[1]],2,2)
table.adult <- data.frame(cols12[order(dimnames(cols12)[[1]]),],
  cols34[order(dimnames(cols34)[[1]]),],
  cols56[order(dimnames(cols56)[[1]]),],
  cols78[order(dimnames(cols78)[[1]]),],
  cols910[order(dimnames(cols910)[[1]]),])
table.adult
##----- check 2way for one target-----
dis1 <- disclosure(synadult,adultsample, target = "workclass",
  keys = c("age" , "occupation","race", "sex" ))
print(dis1, to.print = c("attrib","check\_2way"))

```